Comparing Lexical, Static-Embedding, and Contextual Drift Signals in Brazilian Portuguese Political Discourse

João Victor M. Correa[1] **, Renato Moraes Silva**[2]

[1]Interinstitutional Center for Computational Linguistics (NILC)
Institute of Mathematics and Computer Science
University of São Paulo (USP), São Carlos, Brazil

[2]Department of Computer Science
Federal University of São Carlos (UFSCar)
Sorocaba, São Paulo, Brazil

victormcorrea@usp.br, renatoms@icmc.usp.br

Abstract. *We compare three families of diachronic drift signals in Brazilian Portuguese political discourse: a lexical-profile baseline based on TF-IDF, a static embedding pipeline based on slice-specific Word2Vec models with Orthogonal Procrustes alignment, and a contextual analysis based on a Portuguese BERT checkpoint. The study uses the floor subset of BrPoliCorpus, covering 24 yearly slices from 2000 to 2023 with 428,366 speeches and 63.0 million retained tokens after preprocessing. To make the methods directly comparable, we build a shared panel of 55 lemmas containing method-specific drift candidates, stable controls, and theory seeds. The resulting comparison shows strong disagreement between the lighter baselines, with Spearman correlation of* -0.54 *and zero overlap in their top 15 terms, while contextual BERT remains moderately closer to Word2Vec than to TF-IDF. A filtered contextual panel reduces stable-control leakage and highlights terms such as bloqueio, salário, intervenção, and eleição. Rather than treating one detector as externally validated ground truth, we frame the study as an exploratory comparison of complementary drift sensitivities, interpretability, and computational cost in Brazilian Portuguese political discourse.*

1. Introduction

Political discourse is a productive setting for diachronic NLP because shifts in agenda, coalition structure, and policy framing leave traces in lexical choice and word usage over time. Distributional models can capture some of these movements at scale, but different modeling choices foreground different aspects of change [Hamilton et al. 2016, Kutuzov et al. 2018, Tahmasebi et al. 2021]. In political corpora the situation is especially complex, since a term may drift because of a new policy referent, a rhetorical reframing, a shift in collocational profile, or a reorganization of discourse coalitions—none of which reduces neatly to a single lexical sense change.

That complexity matters for evaluation. Shared-task benchmarks have improved comparability in lexical semantic change detection [Schlechtweg et al. 2020], but curated external labels remain scarce—especially outside the few languages and corpora covered by standard tasks. Distributional pipelines can also react to temporal artifacts, alignment

noise, and other confounds when outputs are read too literally [Dubossarsky et al. 2019]. Without externally validated ground truth, the central question becomes not which detector is "correct" but what each detector appears to capture when applied to the same political timeline.

This question is timely for Brazilian Portuguese. Large-scale political corpora such as BrPoliCorpus [Lima-Lopes 2025] are only now becoming available, yet comparative drift studies that place multiple detection families side by side on the same Portuguese political material remain rare. We use the floor subset of BrPoliCorpus—covering 24 yearly slices of Chamber speeches from 2000 to 2023—and compare a TF-IDF lexical-profile baseline, a slice-specific aligned Word2Vec pipeline, and a contextual BERT stage, all scored on a shared comparison panel of 55 lemmas. The paper asks three concrete questions: (i) how strongly do the methods agree on the same lemmas, (ii) which candidates remain method-specific, and (iii) whether the higher-cost contextual stage adds enough interpretive value to justify its expense. That comparison fills a descriptive gap in Portuguese, where published multi-method drift analyses are still much scarcer than in the better-studied English-centered literature.

2. Related Work

Distributional semantic change detection has evolved from co-occurrence comparisons to aligned static embeddings and then to contextual representations [Kutuzov et al. 2018, Tahmasebi et al. 2021]. Among static approaches, slice-specific Word2Vec models aligned with Orthogonal Procrustes have been shown to recover known historical shifts and reveal statistical regularities of change [Hamilton et al. 2016]. Such embeddings remain attractive because they are efficient and produce interpretable neighborhood movements, but they compress all usages of a word into a single vector, blurring cases in which older and newer usages coexist.

Evaluation remains the harder problem. Standardized shared tasks have improved comparability [Schlechtweg et al. 2020], yet many real-world studies still operate without external semantic ground truth. Apparent change can emerge from temporal referencing, alignment noise, or shifts in genre rather than from semantic drift alone [Dubossarsky et al. 2019]. This concern motivates the use of stable controls and agreement analysis in the present work.

Contextual models offer a partial response to the aggregation problem. BERT-style representations encode token-level usages and can surface variation that static spaces smooth over [Giulianelli et al. 2020, Devlin et al. 2019]. Scalable contextual methods that improve interpretability have also been proposed [Montariol et al. 2021], and pre-trained Portuguese checkpoints such as BERTimbau [Souza et al. 2020] have made contextual analysis feasible for Portuguese NLP. We do not position contextual BERT as a final arbiter; instead, we place it beside TF-IDF and aligned Word2Vec on the same political panel so that agreement, disagreement, and computational tradeoffs can be examined directly.

Diachronic analysis of political language has gained traction through parliamentary corpora. Hofmann et al. [Hofmann et al. 2020] compared lexical usage in Austrian political discourse across time, but comparable multi-method studies on Brazilian Portuguese political speech remain scarce. BrPoliCorpus [Lima-Lopes 2025] now provides the scale

Table 1. Corpus summary for the BrPoliCorpus floor subset, which contains Brazilian Chamber floor speeches. Yearly slice statistics are reported as p25/median/p75 across the 24 yearly files; token and character counts are shown in millions.

Dataset	Period	Slices	Corpus totals	Yearly slices (Q1/Med/Q3)
			428,366 speeches	14,648 / 16,752 / 21,987 speeches
			63.0M tokens	2.00 / 2.62 / 3.45M tokens
			538.5M chars	16.9 / 22.3 / 29.7M chars
BrPoliCorpus floor	2000–2023	24	225,484 lemmas	39,456 / 50,269 / 58,637 lemmas

needed for this kind of controlled comparative work.

3. Corpus and Experimental Setup

The experiments use the floor subset of BrPoliCorpus [Lima-Lopes 2025], which contains Brazilian Chamber floor speeches from 2000 to 2023 organized into 24 yearly slices.

3.1. Preprocessing

Raw speeches are tokenized and tagged with `spaCy` (`pt_core_news_sm`). We retain nouns, verbs, and adjectives, exclude proper nouns, remove stopwords, punctuation, and numeric tokens, and keep only tokens with at least two characters. All text is lowercased with accents preserved. Each token is then mapped to its lemma form, producing the `content_lemma` representation used by all three methods. After preprocessing, the corpus retains 428,366 speeches and 63,036,642 tokens across 225,484 unique lemmas. Table 1 summarizes the corpus totals and the quartile structure of yearly variation.

3.2. Candidate Eligibility

A lemma enters the scored vocabulary only if it appears at least 50 times per slice, in at least 5 documents per slice, and in at least 80% of the 24 slices. Drift and stable-control candidates are further restricted to dominant part-of-speech classes (nouns and adjectives), and a small set of procedural high-frequency terms is excluded. These thresholds reduce the scored vocabulary to 3,221 lemmas for the two cheaper methods.

At the slice level, yearly document counts range from 3,041 speeches in 2000 to 26,868 in 2015, with a median of 16,752 documents per year. Retained token volume ranges from 630,426 in 2000 to 4,171,779 in 2003, with a median of 2,620,480 tokens. Document volume is high throughout the panel, but token volume and vocabulary size vary enough to motivate the frequency and slice-presence controls described above.

4. Methodology

Figure 1 summarizes the workflow. The preprocessed discovery corpus feeds two cheaper baselines; their method-led candidates are then combined with stable controls and theory seeds to create the shared panel used for contextual inspection and agreement analysis.

4.1. TF-IDF Baseline

For each eligible lemma we compute a TF-IDF vector across yearly slices using smoothed inverse-document frequency. Drift is measured as the mean absolute change in TF-IDF weight between adjacent slices. This captures shifts in lexical salience—how central a term becomes to each year's agenda—without relying on distributional neighborhoods.

4.2. Word2Vec Baseline

We train independent skip-gram models per yearly slice (300 dimensions, window 5, 5 negative samples, minimum count 5, 5 epochs, 2 replicates with different seeds) and align them into a common space with Orthogonal Procrustes over the top 20,000 anchor words [Hamilton et al. 2016]. Drift is the mean cosine distance between a lemma's aligned vectors in consecutive slices, averaged over replicates. This captures persistent neighborhood reorganization.

4.3. Contextual BERT

We embed sampled in-context occurrences with a Portuguese BERT checkpoint (BERTimbau Large [Souza et al. 2020]), extracting hidden-state layers -1 and -4. For each lemma and slice, up to 64 occurrences are sampled within a 48-token context window; the contextual drift score is the mean cosine distance between slice-level centroid embeddings in consecutive slices. The agreement layer keeps layer -1 as the preferred ranking and uses layer -4 as a robustness view.

4.4. Shared Comparison Panel

Rather than compare unrestricted vocabularies, we construct a shared panel of 55 lemmas: 15 Word2Vec-led drift candidates, 15 TF-IDF-led drift candidates, 20 stable controls (lemmas ranked lowest by both cheap methods), and 5 theory seeds drawn from prior work on Brazilian political vocabulary (*democracia*, *corrupção*, *reforma*, *economia*, *liberdade*). All three methods observe the same candidate set, so disagreement reflects properties of the signal rather than divergent vocabularies. The agreement layer derives pairwise rank correlations, top-k overlaps, stable-control leakage diagnostics, and a filtered contextual ranking.

Figure 2 visualizes the pairwise agreement structure, while Figure 3 summarizes overlap and rank-distribution statistics.

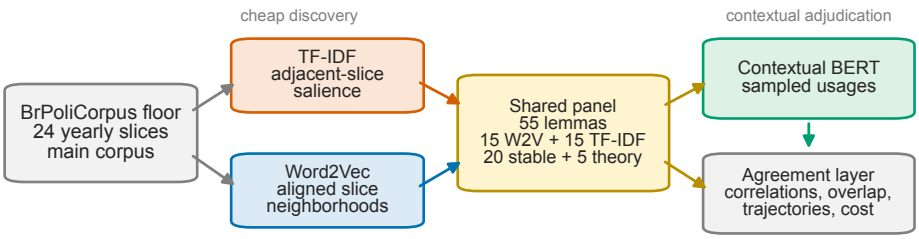

Figure 1. Comparative workflow. The two cheaper baselines nominate candidates; the shared panel combines those candidates with controls and theory seeds before contextual inspection and agreement analysis.

5. Results

The two lighter baselines disagree strongly even after they are constrained to the same candidate universe. On the shared panel, Word2Vec and TF-IDF reach Spearman $\rho = -0.54$ ($p = 2.1 \times 10^{-5}$) and share zero terms in their top 15 rankings. The preferred

contextual layer partly mediates this split without collapsing onto either baseline. Its correlation with Word2Vec is $\rho = 0.21$ ($p = 0.128$) and with TF-IDF $\rho = 0.12$ ($p = 0.365$). These positive associations are directionally informative but not strong enough to justify treating contextual BERT as a substitute for either cheaper method. Figure 2 shows that contextual top ranks draw support from both camps, while the two cheap methods remain sharply separated.

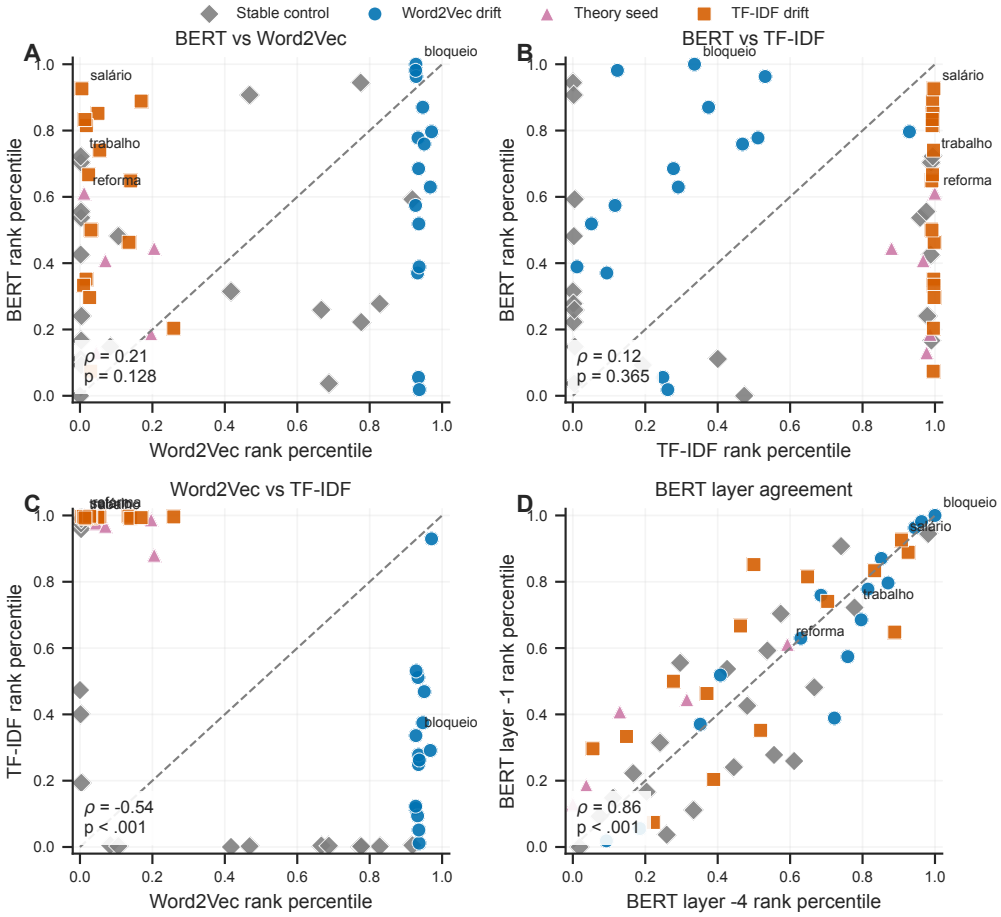

Figure 2. Agreement on the shared panel across TF-IDF, Word2Vec, and the preferred BERT layer; panel (D) compares BERT layers -4 and -1. Annotations report Spearman ρ and two-sided p-values.

Overlap statistics and bucket-level rank distributions prove more informative than raw top-k lists alone. In the preferred BERT layer, the top-15 overlap is 7 terms with Word2Vec and 6 terms with TF-IDF, whereas the baseline-to-baseline overlap remains zero. Figure 3 also shows that selected drift terms are ranked significantly above stable controls by Word2Vec and TF-IDF ($p < .001$ in both cases), while contextual BERT still separates the groups but with weaker margins ($p = .008$). This pattern matches the filtered contextual panel. BERT provides useful adjudication, but the raw contextual ranking still leaks some stable controls and should not be treated as an unrestricted final list. The two contextual extraction layers nevertheless agree strongly with each other (Spearman $= 0.858$), so the broad comparative picture does not depend on a single extraction depth. This high inter-layer agreement is consistent with findings that mid-to-upper layers carry

most of the semantic change signal, while the final layer increasingly encodes surface-form information that can bias change scores [Laicher et al. 2021, Cassotti et al. 2024]; the stability between layers -1 and -4 in our setting suggests that orthographic bias does not yet dominate, though it does not rule it out.

As a robustness check, we recomputed contextual scores using average pairwise distance (APD) [Giulianelli et al. 2020] instead of the centroid-based metric. APD operates on all individual occurrence embeddings rather than on slice-level centroids, giving equal influence to each usage regardless of embedding magnitude. On the preferred layer, APD and centroid rankings correlate at Spearman $\rho = 0.79$ ($p < 10^{-12}$), with 7 of 15 top terms shared. The two metrics foreground different properties: APD promotes high-frequency polysemous terms whose embedding clouds are broad (e.g. *trabalho*, *público*), while the centroid metric favours terms with coherent directional shifts (*bloqueio*, *salário*). Notably, the centroid metric separates method-nominated drift candidates from stable controls more cleanly than APD ($p = .008$ vs. $p = .081$, Mann–Whitney). The strong overall correlation confirms that the comparative conclusions are robust to the choice of contextual aggregation, while the separation difference suggests that centroid distance is a less noisy signal for this panel-based comparison.

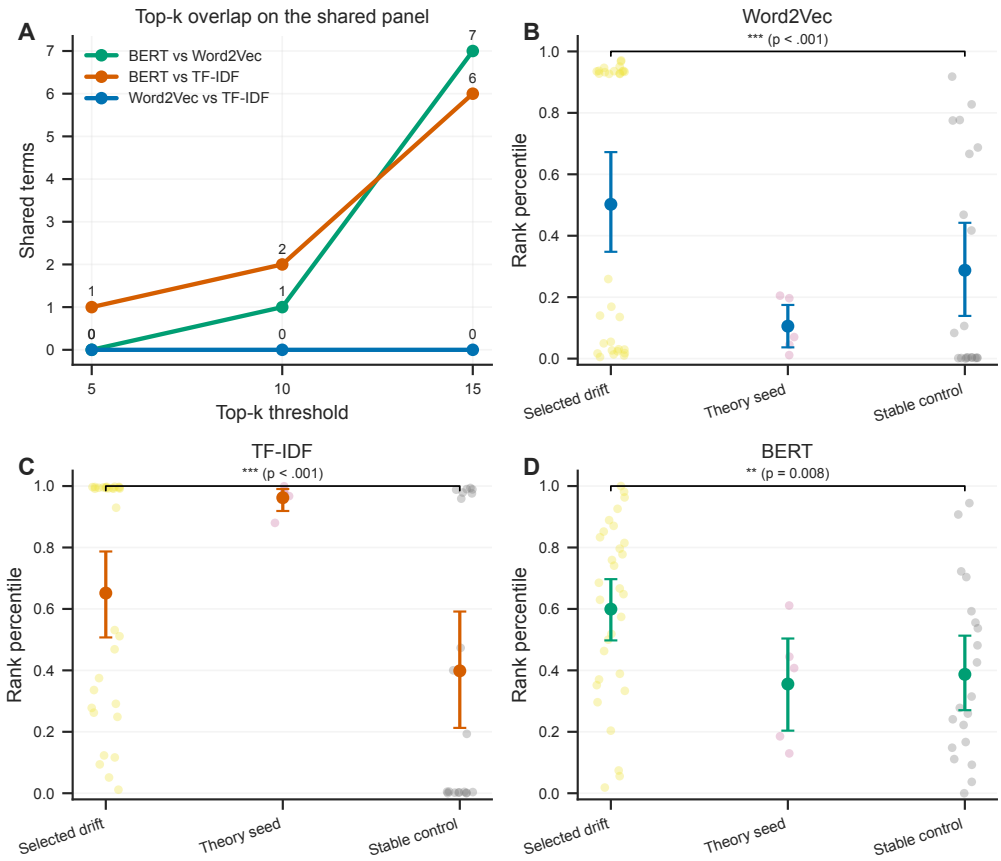

Figure 3. Overlap and rank-distribution summaries on the shared comparison panel. Panel (A) traces top-k overlap; panels (B) through (D) compare rank-percentile distributions for selected drift terms, theory seeds, and stable controls.

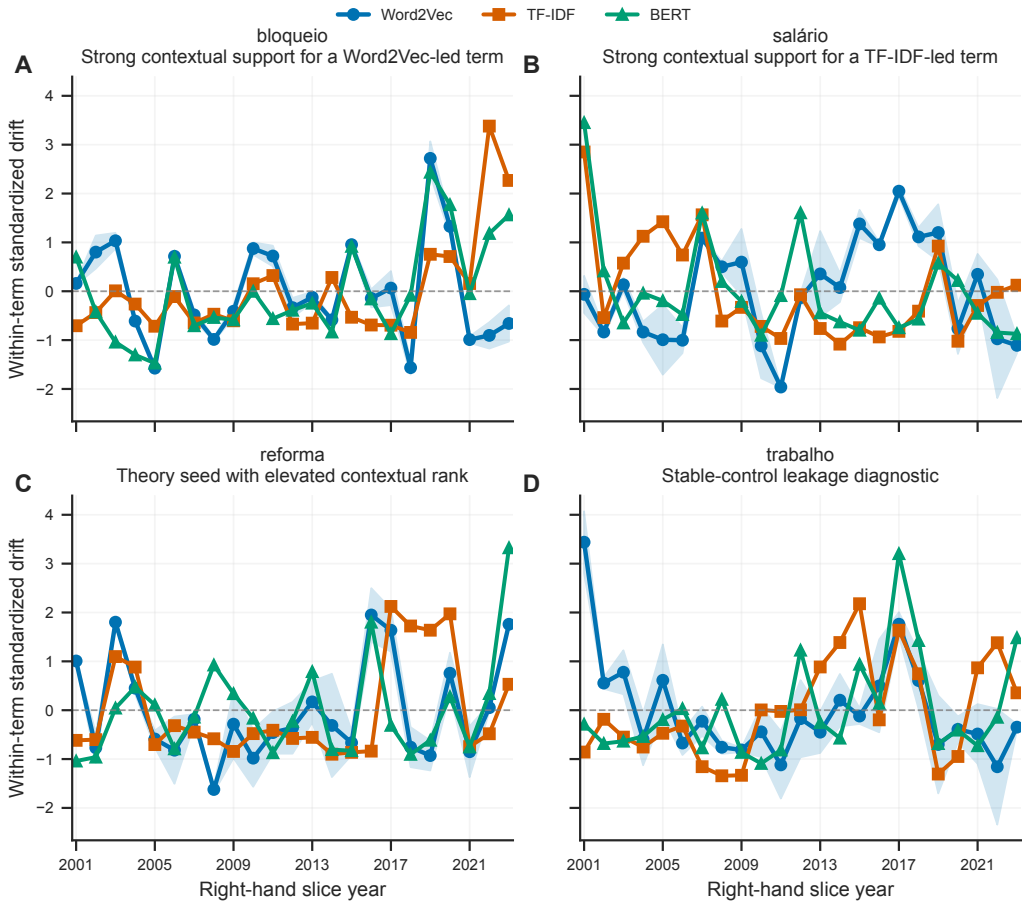

Figure 4. Representative term trajectories across methods. Scores are standardized within term and method to emphasize temporal shape rather than absolute scale.

6. Discussion

Representative term trajectories show that the methods often disagree about *when* a term moves, not only whether it moves at all. Figure 4 illustrates four cases: a Word2Vec-led term with strong contextual support (*bloqueio*), a TF-IDF-led term with strong contextual support (*salário*), a theory seed whose contextual rank rises despite modest static drift (*reforma*), and a stable-control diagnostic term (*trabalho*) that reveals where contextual leakage remains. These cases show why a comparative analysis is more informative than a single-method ranking.

The disagreement itself is substantively informative. TF-IDF appears most sensitive to lexical-profile reweightings. When a term becomes newly central to yearly agendas, policy disputes, or formulaic framing, its salience can move even if its local semantic neighborhood remains comparatively stable. Word2Vec, by contrast, is better positioned to highlight persistent neighborhood reorganization across slices. Contextual BERT sits between these views. Because it observes token-level usages, it can support either a TF-IDF-led or a Word2Vec-led candidate, but the stable-control leakage analysis shows that greater context sensitivity does not automatically produce a cleaner unrestricted ranking.

This interpretation fits the term-level evidence. *Salário* behaves like a lexical-profile

Table 2. Method scope and relative computational cost. The cost column reports order-of-magnitude class and approximate multiplier relative to TF-IDF, measured on the same single-machine run; absolute wall-clock times will vary with hardware.

Method	Signal focus	Scope	Relative cost
TF-IDF	Adjacent-slice salience shift	3,221 lemmas	seconds ($1\times$)
Word2Vec	Aligned neighborhood displacement	3,221 lemmas	minutes ($\sim 56\times$)
Contextual BERT	Usage divergence from sampled contexts	55 lemmas / 82,461 occurrences	hours ($\sim 600\times$)

case with strong contextual support, suggesting that changes in agenda prominence and usage framing are both at work. *Bloqueio* looks more like a neighborhood-displacement case, where the contextual layer aligns with the embedding-led signal. *Reforma* is useful for a different reason: it is not a top filtered contextual term, but as a theory seed it shows that politically central vocabulary can display contextual differentiation even when it is not the strongest candidate in the cheap rankings. The shared panel includes theory seeds and stable controls alongside method-led drift terms for this reason.

The cost comparison in Table 2 sharpens the practical picture. TF-IDF runs in seconds, Word2Vec remains practical at corpus scale in minutes, and contextual BERT is roughly two orders of magnitude more expensive even after being restricted to the 55-lemma panel. Given the modest rank correlations reported above, the current evidence does not support replacing the cheaper methods with BERT. Instead, it supports using contextual modeling selectively as an adjudication layer when disagreement is substantively interesting.

7. Conclusion

The results support an exploratory comparative study rather than an argument for a single best detector of semantic change. TF-IDF and Word2Vec highlight clearly different regions of the candidate space, and contextual BERT adds a useful but higher-cost inspection layer that remains distinct from both. For this Portuguese political setting, disagreement is therefore part of the result, suggesting that the three families are sensitive to different mixtures of lexical salience, neighborhood displacement, and usage-level variation.

8. Limitations

The first limitation is evaluative. This paper does not claim external semantic ground truth for BrPoliCorpus floor, so agreement and disagreement are analyzed as comparative evidence rather than as proof that one detector is correct. The second limitation is methodological. Although the filtered contextual list is useful for interpretation, the raw contextual ranking still admits some stable controls near the top, which means that contextual sensitivity does not automatically yield a cleaner unrestricted candidate list. A third consideration is the contextual aggregation strategy. We report centroid-based drift as the primary contextual metric and verify it against APD [Giulianelli et al. 2020, Cassotti et al. 2024] (Section 5); the strong correlation and the centroid metric's tighter bucket separation support this choice, though APD captures complementary embedding-cloud variance.

Sense-based clustering methods [Montariol et al. 2021] were not attempted here. The fourth limitation is temporal granularity. Yearly slicing is appropriate for a first comparative baseline, but it compresses within-year shocks, campaign cycles, and short-lived institutional disputes that may matter for political language.

Future work could address the evaluative gap through external annotation layers or noisy supervision from labeled parliamentary datasets, though such resources still require careful deduplication and aggregation before they can serve as reliable comparison targets. On the modeling side, adopting task-specialized models such as XL-LEXEME [Cassotti et al. 2024], exploring newer divergence metrics [Rachinskiy and Schlechtweg 2025], and applying orthographic-bias mitigation through subword masking [Laicher et al. 2021] would bring the contextual stage closer to current best practices. A symbolic companion analysis using lexical complexity or readability features could also help separate lexical-semantic drift from rhetorical or stylistic change.

9. Ethics Statement

The study analyzes public institutional political speech and reports only aggregate corpus statistics, ranked lemmas, and usage-level comparative signals. We do not use the models to profile individual speakers or to infer private attributes. Even so, political corpora encode unequal representation, partisan conflict, and occasionally harmful rhetoric, so drift scores should not be read as normative judgments, direct evidence of speaker intent, or standalone claims about social meaning. The outputs in this paper are exploratory comparative signals whose interpretation depends on corpus composition, preprocessing choices, and the distinction between lexical salience, contextual variation, and broader discourse change.

References

Cassotti, P., Ferrara, A., Basile, P., de Gemmis, M., and Semeraro, G. (2024). A systematic comparison of contextualized word embeddings for lexical semantic change. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING)*, pages 2717–2728, Torino, Italy. European Language Resources Association and International Committee on Computational Linguistics.

Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. In Burstein, J., Doran, C., and Solorio, T., editors, *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.

Dubossarsky, H., Hengchen, S., Tahmasebi, N., and Schlechtweg, D. (2019). Time-out: Temporal referencing for robust modeling of lexical semantic change. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 457–470. Association for Computational Linguistics.

Giulianelli, M., Del Tredici, M., and Fernández, R. (2020). Analysing lexical semantic change with contextualised word representations. In Jurafsky, D., Chai, J., Schluter, N.,

and Tetreault, J., editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 3960–3973, Online. Association for Computational Linguistics.

Hamilton, W. L., Leskovec, J., and Jurafsky, D. (2016). Diachronic word embeddings reveal statistical laws of semantic change. In Erk, K. and Smith, N. A., editors, *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1489–1501, Berlin, Germany. Association for Computational Linguistics.

Hofmann, K., Marakasova, A., Baumann, A., Neidhardt, J., and Wissik, T. (2020). Comparing lexical usage in political discourse across diachronic corpora. In *Proceedings of the Second ParlaCLARIN Workshop*, pages 58–65, Marseille, France. European Language Resources Association.

Kutuzov, A., Øvrelid, L., Szymanski, T., and Velldal, E. (2018). Diachronic word embeddings and semantic shifts: a survey. In Bender, E. M., Derczynski, L., and Isabelle, P., editors, *Proceedings of the 27th International Conference on Computational Linguistics*, pages 1384–1397, Santa Fe, New Mexico, USA. Association for Computational Linguistics.

Laicher, S., Kurtyigit, S., Schlechtweg, D., Kuber, J., and Schulte im Walde, S. (2021). Explaining and improving BERT performance on lexical semantic change detection. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Student Research Workshop*, pages 192–202, Online. Association for Computational Linguistics.

Lima-Lopes, R. E. d. (2025). Brpolicorpus: Brazilian political corpus.

Montariol, S., Martinc, M., and Pivovarova, L. (2021). Scalable and interpretable semantic change detection. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4642–4652, Online. Association for Computational Linguistics.

Rachinskiy, A. and Schlechtweg, D. (2025). Rethinking evaluation metrics for lexical semantic change detection. *arXiv preprint arXiv:2602.15716*.

Schlechtweg, D., McGillivray, B., Hengchen, S., Dubossarsky, H., and Tahmasebi, N. (2020). SemEval-2020 task 1: Unsupervised lexical semantic change detection. In Herbelot, A., Zhu, X., Palmer, A., Schneider, N., May, J., and Shutova, E., editors, *Proceedings of the Fourteenth Workshop on Semantic Evaluation*, pages 1–23, Barcelona (online). International Committee for Computational Linguistics.

Souza, F., Nogueira, R., and Lotufo, R. (2020). BERTimbau: Pretrained BERT models for Brazilian Portuguese. In *Intelligent Systems – BRACIS 2020*, pages 403–417, Rio Grande, Brazil. Springer.

Tahmasebi, N., Borin, L., Jatowt, A., Xu, Y., and Hengchen, S., editors (2021). *Computational Approaches to Semantic Change*. Language Science Press, Berlin.